

TEMPDB CONTENTION

HOW TO IDENTIFY AND RESOLVE IT

Haripriya Naidu

SQL DBA

She/Her

About me

- SQL Database Administrator – 10 years
- Focus on
 - SQL Server Engine Internals
 - Performance Tuning
 - Server/DB Operations
- Actively participate in
 - NESQL
 - SQL Saturday Boston
 - Multiple online learning communities
- AWS Certified Solutions Architect
- Rookie of the Year 2023 @Work



Agenda



Uses of TempDB



Types of
Contention



Identifying
Contention



Ways to Resolve

Why learn about TempDB Contention



Troubleshoot
Slowness



Concurrency
Management

What runs on TempDB



The diagram consists of three identical rectangular boxes arranged horizontally. Each box has a dark blue header bar at the top and a light blue body. The text is centered within each box. The first box is labeled 'User Objects', the second 'Internal Objects', and the third 'Version Stores'.

User
Objects

Internal
Objects

Version
Stores

User objects

- Local and Global Temp Tables - #table, ##table
- Local and Global Temp Stored Procedures - #proc, ##proc
- Table Variables - @table
- Appear in sys.all_objects view

Internal objects

- Worktables to store intermediate results for
 - Sorts, Hash Joins, Aggregates
 - Spools, Cursors
 - INSTEAD OF triggers
 - Store Service Broker messages in transit
 - Online Index Rebuild when SORT_IN_TEMPDB is ON
- Users cannot create these internal objects

Version Stores

- Used for Row versioning for the following:
 - Snapshot isolation levels
 - Online index operations
 - AFTER triggers
 - MARS(Multiple active result sets)

| Types of Contention



Object Allocation Contention – DML



Metadata Contention – DDL



Temp Table Cache Contention

Object Allocation Contention - DML

- In tempdb, there are numerous insert, update, delete operations running concurrently on temporary objects, waiting for a latch on PFS, GAM, SGAM pages on a datafile, causing contention

| Datafile – tempdb.mdf

Page 0: File Header

Page 1: Page Free Space (PFS)

Page 2: Global Allocation Map (GAM)

Page 3: Shared Global Allocation Map (SGAM)

Page 4: Unused

Page 5: Unused

Page 6: Differential Change Map (DCM)

Page 7: Bulk Change Map (BCM)

Data pages

... <more data pages>

Tracking pages



- Track the status of data pages
- Used when SQL Server needs to allocate space

Data pages

Page Free Space (PFS)

- Tracks free space available for each page
- All operations need a latch on PFS page to know which page has space

Page 1: PFS

PFS: Page Alloc Status @0x000000F6893F8000

(1:0)	-	(1:3)	=	ALLOCATED	100_PCT_FULL
(1:4)	-	(1:5)	=	NOT ALLOCATED	0_PCT_FULL
(1:6)	-	(1:7)	=	ALLOCATED	100_PCT_FULL
(1:8)	-		=	ALLOCATED	0_PCT_FULL
(1:9)	-		=	ALLOCATED	100_PCT_FULL
(1:10)	-		=	ALLOCATED	0_PCT_FULL
(1:11)	-		=	ALLOCATED	0_PCT_FULL
(1:12)	-		=	ALLOCATED	100_PCT_FULL
(1:13)	-		=	ALLOCATED	0_PCT_FULL
(1:14)	-		=	ALLOCATED	0_PCT_FULL
(1:15)	-		=	ALLOCATED	0_PCT_FULL
(1:16)	-	(1:17)	=	ALLOCATED	0_PCT_FULL
(1:18)	-		=	ALLOCATED	0_PCT_FULL
(1:19)	-	(1:20)	=	ALLOCATED	0_PCT_FULL
(1:21)	-	(1:22)	=	ALLOCATED	0_PCT_FULL
(1:23)	-		=	ALLOCATED	0_PCT_FULL
(1:24)	-	(1:31)	=	ALLOCATED	0_PCT_FULL
(1:32)	-		=	ALLOCATED	50_PCT_FULL
(1:33)	-		=	ALLOCATED	0_PCT_FULL
(1:34)	-		=	ALLOCATED	0_PCT_FULL
(1:35)	-		=	ALLOCATED	0_PCT_FULL
(1:36)	-	(1:38)	=	ALLOCATED	0_PCT_FULL
(1:39)	-		=	ALLOCATED	0_PCT_FULL
(1:40)	-		=	ALLOCATED	0_PCT_FULL
(1:41)	-		=	ALLOCATED	0_PCT_FULL
(1:42)	-	(1:44)	=	ALLOCATED	0_PCT_FULL
(1:45)	-		=	ALLOCATED	0_PCT_FULL
(1:46)	-	(1:48)	=	ALLOCATED	0_PCT_FULL
(1:49)	-		=	ALLOCATED	0_PCT_FULL
(1:50)	-	(1:54)	=	ALLOCATED	0_PCT_FULL

GAM (Global Allocation Map)

Page 2: GAM

- Tracks if an uniform extent is allocated or not
- 1 extent = 8 pages
- Uniform extent = all 8 pages belong to same object
- All operations need a latch on GAM page to know which page is allocated or not

GAM: Extent Alloc Status @0x0000000F6BE1F80C2

(1:0)	-	(1:150608)	=	ALLOCATED
(1:150616)	-		=	NOT ALLOCATED
(1:150624)	-	(1:150656)	=	ALLOCATED
(1:150664)	-		=	NOT ALLOCATED
(1:150672)	-	(1:158568)	=	ALLOCATED
(1:158576)	-		=	NOT ALLOCATED
(1:158584)	-	(1:158808)	=	ALLOCATED
(1:158816)	-		=	NOT ALLOCATED
(1:158824)	-	(1:161768)	=	ALLOCATED
(1:161776)	-		=	NOT ALLOCATED
(1:161784)	-	(1:165088)	=	ALLOCATED
(1:165096)	-		=	NOT ALLOCATED
(1:165104)	-	(1:166984)	=	ALLOCATED
(1:166992)	-		=	NOT ALLOCATED
(1:167000)	-	(1:170720)	=	ALLOCATED
(1:170728)	-		=	NOT ALLOCATED
(1:170736)	-	(1:176984)	=	ALLOCATED
(1:176992)	-		=	NOT ALLOCATED
(1:177000)	-	(1:177224)	=	ALLOCATED
(1:177232)	-		=	NOT ALLOCATED
(1:177240)	-	(1:177352)	=	ALLOCATED
(1:177360)	-		=	NOT ALLOCATED
(1:177368)	-	(1:178808)	=	ALLOCATED
(1:178816)	-		=	NOT ALLOCATED
(1:178824)	-	(1:181368)	=	ALLOCATED
(1:181376)	-		=	NOT ALLOCATED
(1:181384)	-	(1:186312)	=	ALLOCATED
(1:186320)	-		=	NOT ALLOCATED
(1:186328)	-	(1:186720)	=	ALLOCATED
(1:186728)	-		=	NOT ALLOCATED

SGAM (Shared Global Allocation Map)

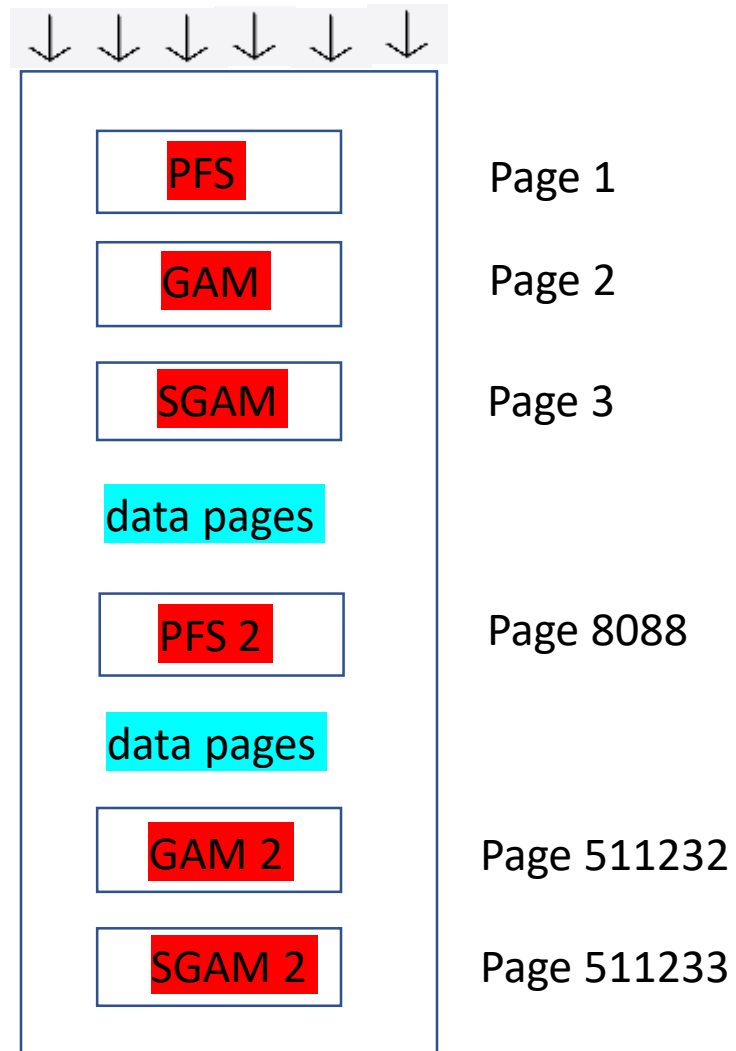
- Tracks if a mixed extent is allocated or not
- Mixed extent = 8 pages belong to different objects
- All operations need a latch on SGAM page to know which page is allocated or not

Page 3: SGAM

SGAM: Extent Alloc Status @0x000000F6B57F80C2

(1:0)	- (1:147648)	= NOT ALLOCATED
(1:147656)	-	= ALLOCATED
(1:147664)	- (1:511224)	= NOT ALLOCATED

Multiple transactions hit on a datafile



Tempdb.mdf

<https://gohigh.substack.com/p/pfs-gam-sgam-pages>

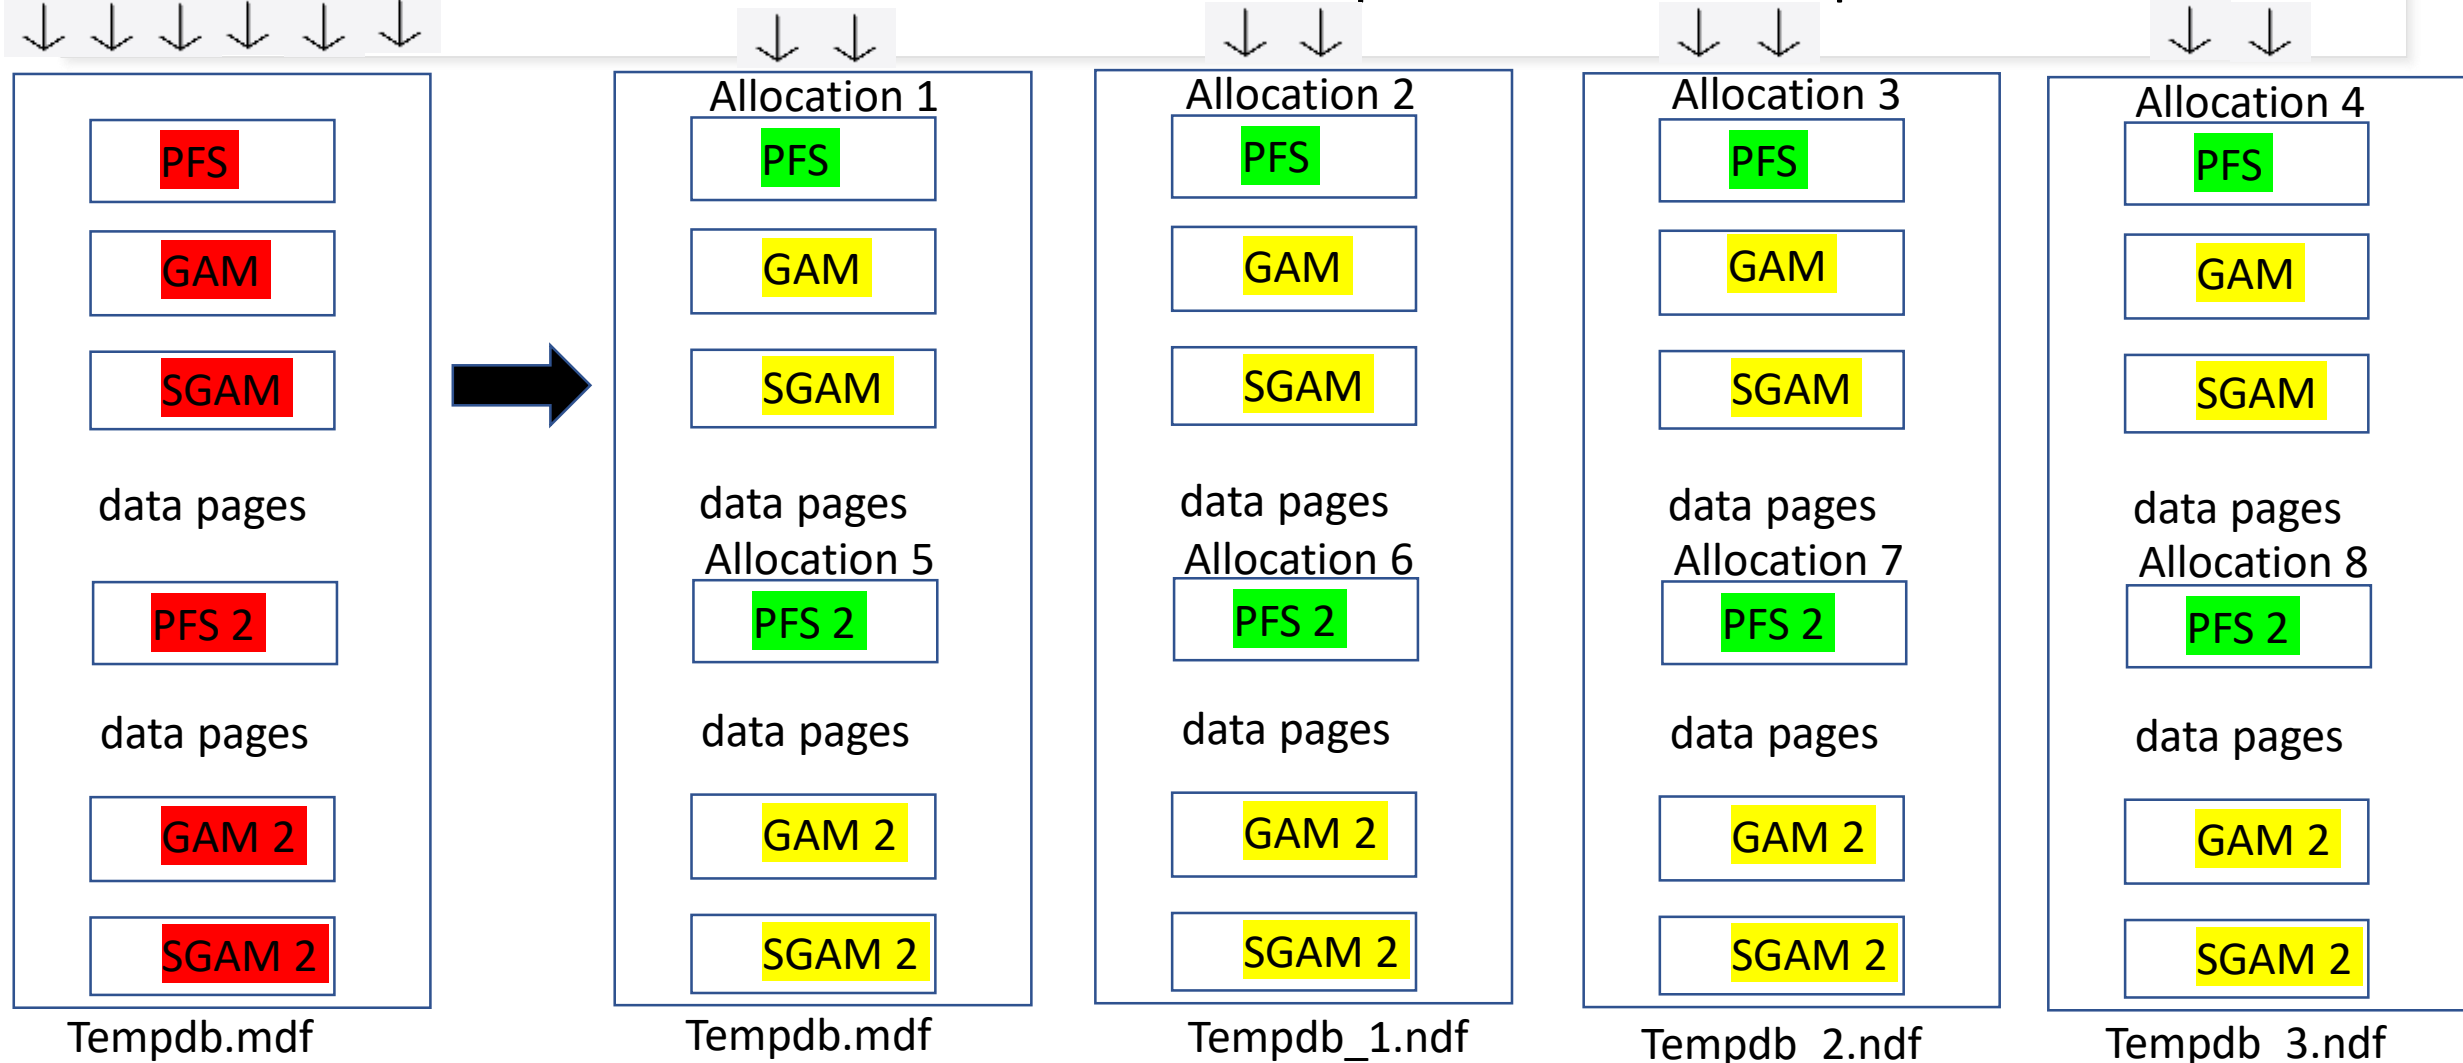
Identify Object Allocation Contention

- PAGELATCH_UP, PAGELATCH_SH waits on PFS and (S)GAM pages
- How to find them:
 - 2:FileID:1 or 2:FileID:<any PFS or (S)GAM page>
- Perfmon counters:
 - Access Methods::Worktables created/sec
 - Access Methods::Workfiles created/sec
 - Number > 200

Resolve Object allocation contention

Multiple Transactions

Transactions are spread across multiple datafiles



Resolve Object allocation contention

- Round Robin
 - between datafiles
 - between PFS pages starting in SQL Server 2017 CU7
- Create 1 data file per logical processor up to 8 logical processors
- Increase additional files in numbers of 4 as needed
- Create with same size and same autogrowth
- Proportionally fills datafiles based on empty space

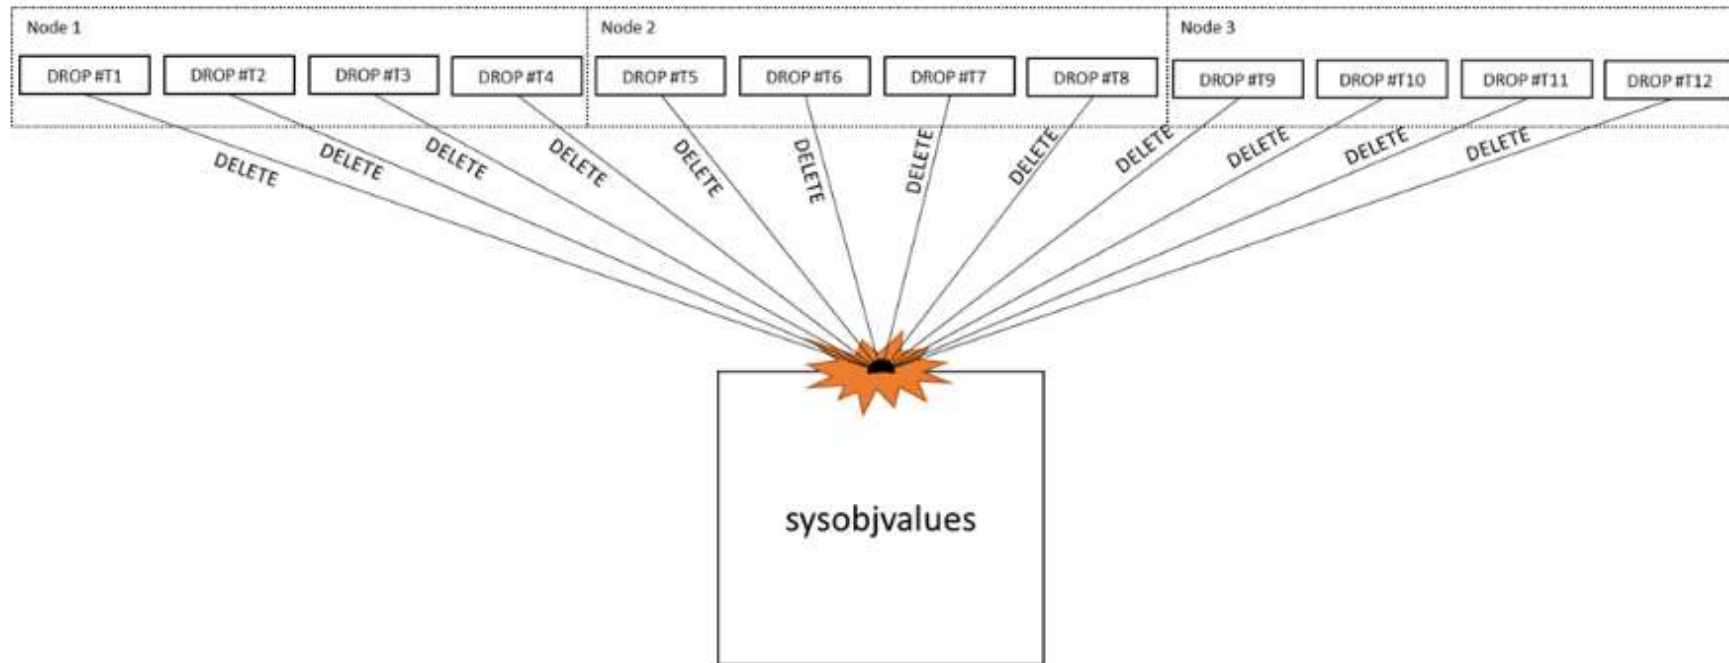
Resolve Object allocation contention

– by SQL Server version

- Prior to SQL 2016: Turn ON Trace flags 1117, 1118
- SQL 2016 and onwards: Trace Flags are ON by default
- SQL 2019: PFS page is updated with a shared latch instead of exclusive latch
- SQL 2022: GAM and SGAM pages are updated with a shared latch
- Follow best practices across all versions

Metadata contention – DDL

- In tempdb, there are numerous create, alter, drop sessions running concurrently on temporary objects, waiting for a latch on system object pages, causing contention



Identify Metadata Contention

- PAGELATCH_EX, PAGELATCH_SH waits on system catalog tables - such as sysobjvalues and sysschobjs
- How to find them 2:FileID:<system page>
- Perfmon counters:
 - SQLStatistics::PageLatch Wait
 - Temp Tables Creation Rate
 - Temp Tables Destruction Rate

Resolve Metadata contention

- Don't drop/truncate temp tables
- Cache temporary objects:
 - Don't alter temp table DDL statements after creation
 - Create indexes inline with table definition
 - Create temp objects within SP, trigger, function
- Enable Memory-Optimized TempDB Metadata starting SQL 2019
 - CU2 moves sysallocunits, sysrowsets
 - Not ON by default

Temp table cache contention

- Contention on cache for Temp table memory objects
- Why?
 - Inefficient usage of Hash tables in cache

Identify/Resolve Temp table cache contention

- CMEMTHREAD waits
- SOS_CACHESTORE spinlock waits
- Open a Microsoft case, needs advanced debugging

DEMO



Summary: Object Allocation Contention

- **Observe:** PFS, GAM, SGAM pages
- **Wait Type:** PAGELATCH_UP, PAGELATCH_SH
- **Resolution:**
 - Multiple Datafiles
 - Upgrade to SQL Server Version 2022

Summary: Metadata Contention

- **Observe:** System Object pages
- **Wait Type:** PAGELATCH_EX, PAGELATCH_SH
- **Resolution:**
 - Don't drop Temp objects
 - Cache Temp objects
 - Turn ON in-memory TempDB metadata feature in SQL Server Version 2019 and above

Summary: Temp Table Cache Contention

- **Observe:** Huge memory objects in cache
- **Wait Type:** CMEMTHREAD, SOS_CACHESTORE
- **Resolution:**



Why learn about TempDB Contention



Troubleshoot
Slowness



Concurrency
Management

Resources

- TempDB: The Good, The Bad, and The Ugly by Pam Lahoud
 - <https://youtu.be/5jpfy19yl7k?si=vmgf3zqxokurpeo0>
- OStress utility: <https://learn.microsoft.com/en-us/troubleshoot/sql/tools/replay-markup-language-utility>
- Klaus Aschenbrenner:
<https://www.sqlservercentral.com/blogs/improved-temp-table-caching-in-sql-server-2014>
- Pro SQL Server Internals book by Dmitri Korotkevitch

Q & A!

Reach Out For More Questions

- Email: haripriya.naidu1@gmail.com
- LinkedIn: <https://www.linkedin.com/in/haripriya-naidu1215>
- GitHub: <https://github.com/haripriyasb/tempdb>
- Substack: <https://gohigh.substack.com>

Thank You!
